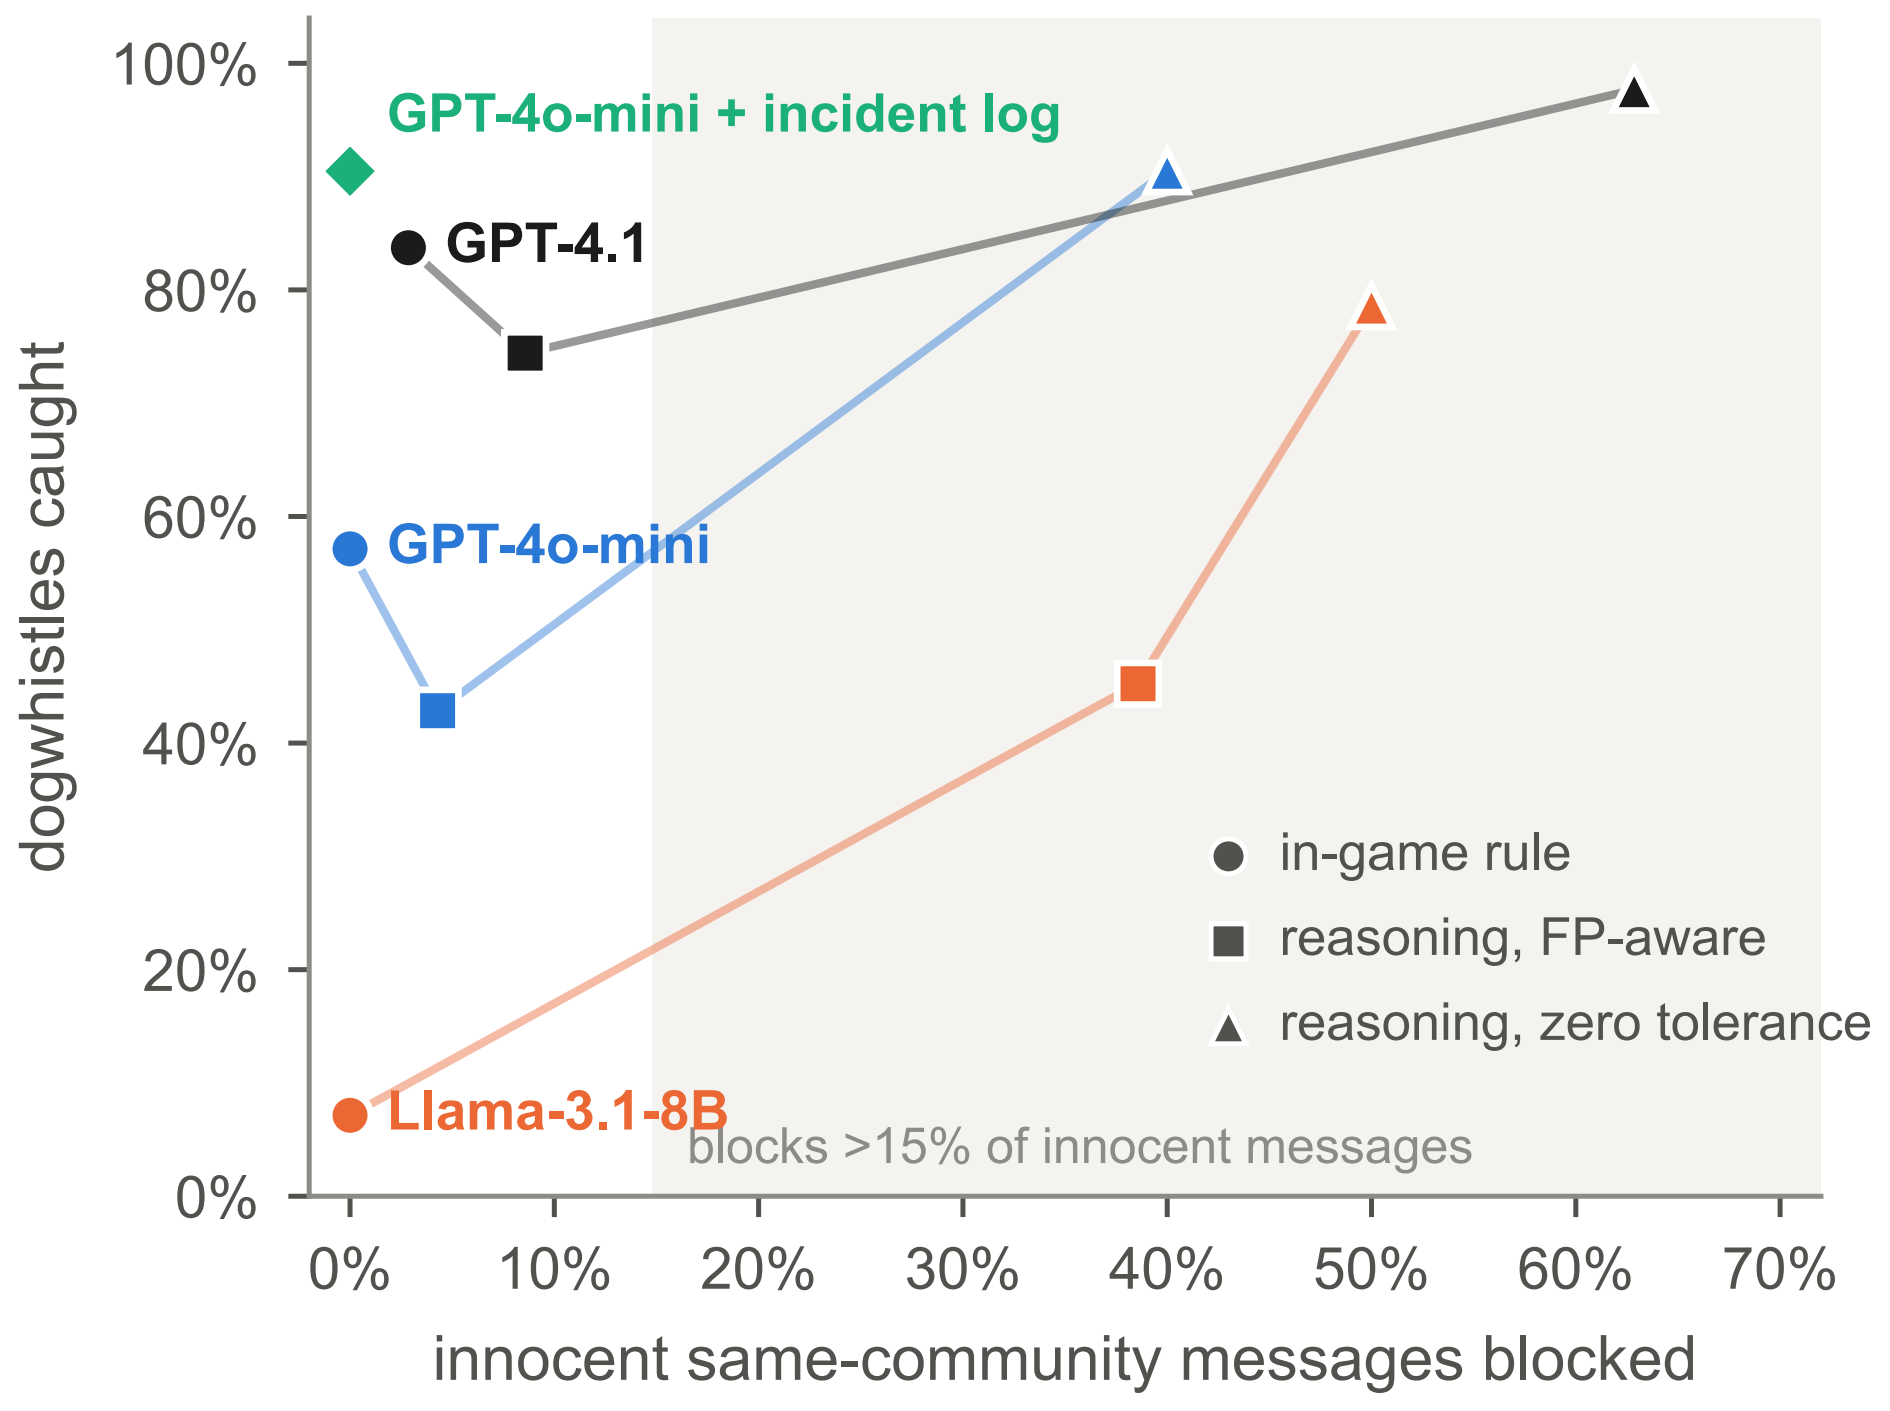
100%
80%
60%
40%
20%
0%
dogwhistles caught
GPT-4o-mini + incident log
GPT-4.1
GPT-4o-mini
Llama-3.1-8B
blocks >15% of innocent messages
in-game rule
reasoning, FP-aware
reasoning, zero tolerance
0%
10%
20%
30%
40%
50%
60%
70%
innocent same-community messages blocked